

Vergleich verschiedener Deep-Learning-Modelle zur Klassifizierung von Fledermausrufen

Christian Müller
Gutenbergstraße 20
64342 Seeheim-Jugenheim
chrmue44@gmail.com

1 Einleitung

Ziel der Untersuchung ist es, einen automatischen Klassifizierer für Fledermausrufe zu entwickeln. Die Klassifizierung soll eine möglichst gute Vorhersage-Genauigkeit bei möglichst geringem Rechenaufwand liefern. Als Software-Plattform dienen das AI-Framework tensorflow 2.0ⁱ und diverse weitere Python-Bibliotheken. Die Klassifizierung in dieser Untersuchung beschränkt sich auf 9 Fledermausarten.

Als Testdaten wurden Aufnahmen vom Naturkundemuseum Berlinⁱⁱ, von Rudolf Böhmⁱⁱⁱ und Christian Müller verwendet. Zum Vergleich verschiedener Deep-Learning Modell-Architekturen wurde immer der gleiche Testdatensatz verwendet.

Inhaltsverzeichnis

1 Einleitung.....	1
2 Verwendete Testdaten.....	2
3 Aufbereitung der Daten.....	4
4 Ergebnisse.....	5
5 Implementierung der untersuchten Modelle.....	6
5.1 Rnn1aModel.....	6
5.2 Rnn5Model.....	7
5.3 rnn6Model.....	8
5.4 rnn6aModel.....	9
5.5 ResNet34.....	10
6 Quellen.....	13

2 Verwendete Testdaten

Folgende Fledermausarten bzw. Klassen sollen detektiert werden:

Fledermausart	Quelle	Anzahl Rufe
Mopsfledermaus (<i>Barbastella barbastellus</i>)	Rudolf Böhm	748
Breitflügelfledermaus (<i>Eptesicus serotinus</i>)	Tierstimmenarchiv.de	1069
Wasserfledermaus (<i>Myotis daubentonii</i>)	Tierstimmenarchiv.de	3308
Großes Mausohr (<i>Myotis myotis</i>)	Tierstimmenarchiv.de	999
Kleinabendsegler (<i>Nyctalus leisleri</i>)	Tierstimmenarchiv.de	779
Großer Abendsegler (<i>Nyctalus noctula</i>)	Tierstimmenarchiv.de	852
Rauhautfledermaus (<i>Pipistrellus nathusii</i>)	Tierstimmenarchiv.de	1678
Zwergfledermaus (<i>Pipistrellus pipistrellus</i>)	Tierstimmenarchiv.de	3579
Mückenfledermaus (<i>Pipistrellus pygmaeus</i>)	Rudolf, Böhm	4450
Heuschrecken (nachfolgend als <i>Cric</i> bezeichnet)	Christian Müller,	424
	Christian Müller,	
Sonstiges (incl. diverser Störgeräusche)	www.github.org/paddygoat ^{iv}	4329

Die Testdaten werden in 3 Klassen aufgeteilt:

Name	Anzahl Rufe	Beschreibung
train	17676	Daten, mit denen das Modell trainiert wird
Dev	2209	Daten, mit denen das Modell während des Trainings überprüft wird
test	2210	Daten, mit denen das Modell am Ende des Trainings validiert wird

3 **Aufbereitung der Daten**

Bevor die Tonaufnahmen dem Training bzw. der Klassifizierung zugeführt werden, erfolgt eine Aufbereitung der Daten. Dazu werden die folgenden Schritte durchgeführt:

1. Resampling der WAV-Aufnahmen zu einer Abtastrate von 312500 Hz
2. Erkennung und Isolation von einzelnen Rufen in der jeweiligen Aufnahme. Hierzu wird das biocoustics package^v verwendet.
3. FFT: Die Aufnahme wird per FFT in den Frequenzbereich transformiert
4. Normierung: Die lauteste festgestellte Amplitude innerhalb eines Rufes wird auf den Wert 1.0 gesetzt. Sämtliche anderen Amplituden werden zu dieser Maximalamplitude ins Verhältnis gesetzt. Mit dieser Maßnahme werden die teilweise erheblichen Lautstärkeunterschiede zwischen den Aufnahmen ausgeglichen.
5. Die FFTs der einzelnen Rufe werden auf ein Format von 129 x 330 Punkten begrenzt. Kürzere Rufe werden mit 0 aufgefüllt, längere Rufe werden am Ende abgeschnitten.
6. Rauschunterdrückung: 2 verschiedene Methoden getestet:
 - alle Amplituden im Spektrogramm unterhalb eines Schwellwertes auf 0 gesetzt
 - nur die oberen 95% der Energie einer Spektrallinie werden berücksichtigt, der Rest auf 0 gesetzt (Methode 1 funktioniert besser)

4 Ergebnisse

In der unten stehenden Tabelle sind die Ergebnisse für die einzelnen Arten aufgelistet. Berechnet ist jeweils die „precision“^{vi} für jede erkannte Klasse bzw. Fledermausart. Generell schneidet das resNet34 am besten ab. Für einzelne Arten liefern andere Modelle jedoch teilweise bessere Ergebnisse.

Das Training der Modelle dauert bei allen Modellen außer dem resNet34 60 ... 90 Sekunden pro Epoche. Lediglich das Resnet34 braucht ca. 1,5 Stunden zum Training einer Epoche, ist also wesentlich ressourcenhungriger.

Modell Art	Precision (TP / (TP + FP)				
	Resnet34	rnn5	rnn6a	flat	rnn1a
----	99.1%	98.8%	99.1%	98.6%	99.1%
Bbar	96.4%	95.2%	91.7%	88.1%	94.0%
Cric	97.5%	97.3%	97.3%	94.7%	100.0%
Eser	80.4%	81.4%	81.4%	76.2%	81.4%
Mdau	98.4%	96.6%	96.9%	91.1%	92.7%
Mmyo	93.5%	80.4%	76.0%	62.0%	64.1%
Nlei	66.7%	74.7%	76%	56.0%	65.3%
Nnoc	78.4%	79.5%	76.1%	76.1%	84.1%
Pnat	96.3%	92.1%	93.3%	91.4%	92.1%
Ppip	98.5%	96.2%	95.0%	94.4%	96.8%
Ppyg	99.8%	99.1%	98.6%	99.5%	99.8%

Da die Modelle zwischen 2 und über 20 Millionen Parameter haben, erscheinen die gut 20.000 Trainingsbeispiele etwas knapp bemessen. Hier würden mehr verlässliche Daten sicher noch zu einer Verbesserung des Ergebnisses beitragen.

5 Implementierung der untersuchten Modelle

5.1 FlatModel

Learning rate: 0.00002, accuracy: 0.91


Model: "flatModel"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 42570)	0
dense (Dense)	(None, 256)	10898176
dense_1 (Dense)	(None, 256)	65792
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 11)	1419

Total params: 10,998,283

Trainable params: 10,998,283

Non-trainable params: 0



5.2 Rnn1aModel

Learning rate: 0.00002, Optimizer „Adam“, 95 Epochs, dropout 0,5

accuracy 0.93


Model: "rnn1aModel"

Layer (type)	Output Shape	Param #
gru (GRU)	(None, 129, 128)	176640
flatten (Flatten)	(None, 16512)	0
dropout (Dropout)	(None, 16512)	0
dense (Dense)	(None, 128)	2113664
dense_1 (Dense)	(None, 11)	1419

Total params: 2,291,723

Trainable params: 2,291,723

Non-trainable params: 0



5.3 Rnn5Model

Learning rate: 0.00002, Optimizer „Adam“, 41 Epochs, dropout 0,5

accuracy 0.943


Model: "rnn5Model"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 29, 196)	970396
batch_normalization (BatchNormalization)	(None, 29, 196)	784
activation (Activation)	(None, 29, 196)	0
dropout (Dropout)	(None, 29, 196)	0
gru (GRU)	(None, 29, 128)	125184
dropout_1 (Dropout)	(None, 29, 128)	0
batch_normalization_1 (BatchNormalization)	(None, 29, 128)	512
gru_1 (GRU)	(None, 29, 128)	99072
dropout_2 (Dropout)	(None, 29, 128)	0
flatten (Flatten)	(None, 3712)	0
dense (Dense)	(None, 128)	475264
dense_1 (Dense)	(None, 11)	1419

Total params: 1,672,631


Trainable params: 1,671,983

Non-trainable params: 648



5.4 rnn6aModel

Learning rate: 0.00002, Epochs: 45, dropout 0.5, accuracy 93.8%



Model: "rnn6aModel"

Layer (type)	Output Shape	Param #
=====		
conv1d (Conv1D)	(None, 29, 196)	970396
batch_normalization (BatchNormalization)	(None, 29, 196)	784
activation (Activation)	(None, 29, 196)	0
dropout (Dropout)	(None, 29, 196)	0
gru (GRU)	(None, 29, 128)	125184
dropout_1 (Dropout)	(None, 29, 128)	0
batch_normalization_1 (BatchNormalization)	(None, 29, 128)	512
gru_1 (GRU)	(None, 29, 128)	99072
dropout_2 (Dropout)	(None, 29, 128)	0
batch_normalization_2 (BatchNormalization)	(None, 29, 128)	512
gru_2 (GRU)	(None, 29, 128)	99072
dropout_3 (Dropout)	(None, 29, 128)	0
flatten (Flatten)	(None, 3712)	0
dense (Dense)	(None, 128)	475264
dense_1 (Dense)	(None, 11)	1419
=====		

Total params: 1,772,215


Trainable params: 1,771,311

Non-trainable params: 904

5.5 ResNet34

rows, timeSteps, classes: 129 330 11, 5 epochs, learning rate 0.00002

accuracy 0.956



Model: "resNet34Model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 129, 330)]	0	[]
reshape (Reshape)	(None, 129, 330, 1)	0	['input_1[0][0]']
zero_padding2d (ZeroPadding2D)	(None, 135, 336, 1)	0	['reshape[0][0]']
conv2d (Conv2D)	(None, 68, 168, 64)	3200	['zero_padding2d[0][0]']
batch_normalization (BatchNormalization)	(None, 68, 168, 64)	256	['conv2d[0][0]']
activation (Activation)	(None, 68, 168, 64)	0	['batch_normalization[0][0]']
max_pooling2d (MaxPooling2D)	(None, 34, 84, 64)	0	['activation[0][0]']
conv2d_1 (Conv2D)	(None, 34, 84, 64)	36928	['max_pooling2d[0][0]']
batch_normalization_1 (BatchNormalization)	(None, 34, 84, 64)	256	['conv2d_1[0][0]']
activation_1 (Activation)	(None, 34, 84, 64)	0	['batch_normalization_1[0][0]']
conv2d_2 (Conv2D)	(None, 34, 84, 64)	36928	['activation_1[0][0]']
batch_normalization_2 (BatchNormalization)	(None, 34, 84, 64)	256	['conv2d_2[0][0]']
add (Add)	(None, 34, 84, 64)	0	['batch_normalization_2[0][0]', 'max_pooling2d[0][0]']
activation_2 (Activation)	(None, 34, 84, 64)	0	['add[0][0]']
conv2d_3 (Conv2D)	(None, 34, 84, 64)	36928	['activation_2[0][0]']
batch_normalization_3 (BatchNormalization)	(None, 34, 84, 64)	256	['conv2d_3[0][0]']
activation_3 (Activation)	(None, 34, 84, 64)	0	['batch_normalization_3[0][0]']
conv2d_4 (Conv2D)	(None, 34, 84, 64)	36928	['activation_3[0][0]']
batch_normalization_4 (BatchNormalization)	(None, 34, 84, 64)	256	['conv2d_4[0][0]']
add_1 (Add)	(None, 34, 84, 64)	0	['batch_normalization_4[0][0]', 'activation_2[0][0]']
activation_4 (Activation)	(None, 34, 84, 64)	0	['add_1[0][0]']
conv2d_5 (Conv2D)	(None, 34, 84, 64)	36928	['activation_4[0][0]']
batch_normalization_5 (BatchNormalization)	(None, 34, 84, 64)	256	['conv2d_5[0][0]']
activation_5 (Activation)	(None, 34, 84, 64)	0	['batch_normalization_5[0][0]']
conv2d_6 (Conv2D)	(None, 34, 84, 64)	36928	['activation_5[0][0]']
batch_normalization_6 (BatchNormalization)	(None, 34, 84, 64)	256	['conv2d_6[0][0]']

```

add_2 (Add)          (None, 34, 84, 64) 0 ['batch_normalization_6[0][0]',
                                           'activation_4[0][0]']
activation_6 (Activation) (None, 34, 84, 64) 0 ['add_2[0][0]']
conv2d_7 (Conv2D)      (None, 17, 42, 128) 73856 ['activation_6[0][0]']
batch_normalization_7 (BatchNormalization) (None, 17, 42, 128) 512 ['conv2d_7[0][0]']
activation_7 (Activation) (None, 17, 42, 128) 0 ['batch_normalization_7[0][0]']
conv2d_8 (Conv2D)      (None, 17, 42, 128) 147584 ['activation_7[0][0]']
batch_normalization_8 (BatchNormalization) (None, 17, 42, 128) 512 ['conv2d_8[0][0]']
conv2d_9 (Conv2D)      (None, 17, 42, 128) 8320 ['activation_6[0][0]']
add_3 (Add)           (None, 17, 42, 128) 0 ['batch_normalization_8[0][0]',
                                           'conv2d_9[0][0]']
activation_8 (Activation) (None, 17, 42, 128) 0 ['add_3[0][0]']
conv2d_10 (Conv2D)     (None, 17, 42, 128) 147584 ['activation_8[0][0]']
batch_normalization_9 (BatchNormalization) (None, 17, 42, 128) 512 ['conv2d_10[0][0]']
activation_9 (Activation) (None, 17, 42, 128) 0 ['batch_normalization_9[0][0]']
conv2d_11 (Conv2D)     (None, 17, 42, 128) 147584 ['activation_9[0][0]']
batch_normalization_10 (BatchNormalization) (None, 17, 42, 128) 512 ['conv2d_11[0][0]']
add_4 (Add)           (None, 17, 42, 128) 0 ['batch_normalization_10[0][0]',
                                           'activation_8[0][0]']
activation_10 (Activation) (None, 17, 42, 128) 0 ['add_4[0][0]']
conv2d_12 (Conv2D)     (None, 17, 42, 128) 147584 ['activation_10[0][0]']
batch_normalization_11 (BatchN (None, 17, 42, 128) 512 ['conv2d_12[0][0]']
ormalization)
activation_11 (Activation) (None, 17, 42, 128) 0 ['batch_normalization_11[0][0]']
conv2d_13 (Conv2D)     (None, 17, 42, 128) 147584 ['activation_11[0][0]']
batch_normalization_12 (BatchNormalization) (None, 17, 42, 128) 512 ['conv2d_13[0][0]']
add_5 (Add)           (None, 17, 42, 128) 0 ['batch_normalization_12[0][0]',
                                           'activation_10[0][0]']
activation_12 (Activation) (None, 17, 42, 128) 0 ['add_5[0][0]']
conv2d_14 (Conv2D)     (None, 17, 42, 128) 147584 ['activation_12[0][0]']
batch_normalization_13 (BatchN ormalization) (None, 17, 42, 128) 512 ['conv2d_14[0][0]']
activation_13 (Activation) (None, 17, 42, 128) 0 ['batch_normalization_13[0][0]']
conv2d_15 (Conv2D)     (None, 17, 42, 128) 147584 ['activation_13[0][0]']
batch_normalization_14 (BatchNormalization) (None, 17, 42, 128) 512 ['conv2d_15[0][0]']
add_6 (Add)           (None, 17, 42, 128) 0 ['batch_normalization_14[0][0]',
                                           'activation_12[0][0]']
activation_14 (Activation) (None, 17, 42, 128) 0 ['add_6[0][0]']
conv2d_16 (Conv2D)     (None, 9, 21, 256) 295168 ['activation_14[0][0]']
batch_normalization_15 (BatchN ormalization) (None, 9, 21, 256) 1024 ['conv2d_16[0][0]']
activation_15 (Activation) (None, 9, 21, 256) 0 ['batch_normalization_15[0][0]']
conv2d_17 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_15[0][0]']
batch_normalization_16 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_17[0][0]']
conv2d_18 (Conv2D)     (None, 9, 21, 256) 33024 ['activation_14[0][0]']
add_7 (Add)           (None, 9, 21, 256) 0 ['batch_normalization_16[0][0]',
                                           'conv2d_18[0][0]']
activation_16 (Activation) (None, 9, 21, 256) 0 ['add_7[0][0]']
conv2d_19 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_16[0][0]']
batch_normalization_17 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_19[0][0]']
activation_17 (Activation) (None, 9, 21, 256) 0 ['batch_normalization_17[0][0]']
conv2d_20 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_17[0][0]']
batch_normalization_18 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_20[0][0]']
add_8 (Add)           (None, 9, 21, 256) 0 ['batch_normalization_18[0][0]',
                                           'activation_16[0][0]']
activation_18 (Activation) (None, 9, 21, 256) 0 ['add_8[0][0]']
conv2d_21 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_18[0][0]']
batch_normalization_19 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_21[0][0]']
activation_19 (Activation) (None, 9, 21, 256) 0 ['batch_normalization_19[0][0]']
conv2d_22 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_19[0][0]']
batch_normalization_20 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_22[0][0]']
add_9 (Add)           (None, 9, 21, 256) 0 ['batch_normalization_20[0][0]',
                                           'activation_18[0][0]']
activation_20 (Activation) (None, 9, 21, 256) 0 ['add_9[0][0]']
conv2d_23 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_20[0][0]']
batch_normalization_21 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_23[0][0]']
activation_21 (Activation) (None, 9, 21, 256) 0 ['batch_normalization_21[0][0]']
conv2d_24 (Conv2D)     (None, 9, 21, 256) 590080 ['activation_21[0][0]']
batch_normalization_22 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_24[0][0]']
add_10 (Add)          (None, 9, 21, 256) 0 ['batch_normalization_22[0][0]',

```

```

activation_22 (Activation) (None, 9, 21, 256) 0 ['activation_20[0][0]']
conv2d_25 (Conv2D) (None, 9, 21, 256) 590080 ['add_10[0][0]']
batch_normalization_23 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_25[0][0]']
activation_23 (Activation) (None, 9, 21, 256) 0 ['batch_normalization_23[0][0]']
conv2d_26 (Conv2D) (None, 9, 21, 256) 590080 ['activation_23[0][0]']
batch_normalization_24 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_26[0][0]']
add_11 (Add) (None, 9, 21, 256) 0 ['batch_normalization_24[0][0]',
'activation_22[0][0]']
activation_24 (Activation) (None, 9, 21, 256) 0 ['add_11[0][0]']
conv2d_27 (Conv2D) (None, 9, 21, 256) 590080 ['activation_24[0][0]']
batch_normalization_25 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_27[0][0]']
activation_25 (Activation) (None, 9, 21, 256) 0 ['batch_normalization_25[0][0]']
conv2d_28 (Conv2D) (None, 9, 21, 256) 590080 ['activation_25[0][0]']
batch_normalization_26 (BatchNormalization) (None, 9, 21, 256) 1024 ['conv2d_28[0][0]']
add_12 (Add) (None, 9, 21, 256) 0 ['batch_normalization_26[0][0]',
'activation_24[0][0]']
activation_26 (Activation) (None, 9, 21, 256) 0 ['add_12[0][0]']
conv2d_29 (Conv2D) (None, 5, 11, 512) 1180160 ['activation_26[0][0]']
batch_normalization_27 (BatchNormalization) (None, 5, 11, 512) 2048 ['conv2d_29[0][0]']
activation_27 (Activation) (None, 5, 11, 512) 0 ['batch_normalization_27[0][0]']

conv2d_30 (Conv2D) (None, 5, 11, 512) 2359808 ['activation_27[0][0]']
batch_normalization_28 (BatchNormalization) (None, 5, 11, 512) 2048 ['conv2d_30[0][0]']
conv2d_31 (Conv2D) (None, 5, 11, 512) 131584 ['activation_26[0][0]']
add_13 (Add) (None, 5, 11, 512) 0 ['batch_normalization_28[0][0]',
'conv2d_31[0][0]']
activation_28 (Activation) (None, 5, 11, 512) 0 ['add_13[0][0]']
conv2d_32 (Conv2D) (None, 5, 11, 512) 2359808 ['activation_28[0][0]']
batch_normalization_29 (BatchNormalization) (None, 5, 11, 512) 2048 ['conv2d_32[0][0]']
activation_29 (Activation) (None, 5, 11, 512) 0 ['batch_normalization_29[0][0]']
conv2d_33 (Conv2D) (None, 5, 11, 512) 2359808 ['activation_29[0][0]']
batch_normalization_30 (BatchNormalization) (None, 5, 11, 512) 2048 ['conv2d_33[0][0]']
add_14 (Add) (None, 5, 11, 512) 0 ['batch_normalization_30[0][0]',
'activation_28[0][0]']
activation_30 (Activation) (None, 5, 11, 512) 0 ['add_14[0][0]']
conv2d_34 (Conv2D) (None, 5, 11, 512) 2359808 ['activation_30[0][0]']
batch_normalization_31 (BatchNormalization) (None, 5, 11, 512) 2048 ['conv2d_34[0][0]']
activation_31 (Activation) (None, 5, 11, 512) 0 ['batch_normalization_31[0][0]']
conv2d_35 (Conv2D) (None, 5, 11, 512) 2359808 ['activation_31[0][0]']
batch_normalization_32 (BatchNormalization) (None, 5, 11, 512) 2048 ['conv2d_35[0][0]']
add_15 (Add) (None, 5, 11, 512) 0 ['batch_normalization_32[0][0]',
'activation_30[0][0]']
activation_32 (Activation) (None, 5, 11, 512) 0 ['add_15[0][0]']
average_pooling2d (AveragePooling2D) (None, 3, 6, 512) 0 ['activation_32[0][0]']
flatten (Flatten) (None, 9216) 0 ['average_pooling2d[0][0]']
dense (Dense) (None, 512) 4719104 ['flatten[0][0]']
dense_1 (Dense) (None, 11) 5643 ['dense[0][0]']

```

=====

Total params: 26,025,099

Trainable params: 26,009,867

Non-trainable params: 15,232

6 Quellen

www.tierstimmenarchiv.de

<https://cran.r-project.org/web/packages/bioacoustics/index.html>

<https://www.tensorflow.org/>

<https://www.coursera.org/learn/neural-networks-deep-learning/home/>

<https://www.analyticsvidhya.com/blog/2021/08/how-to-code-your-resnet-from-scratch-in-tensorflow/>

Gulli Anonio, Kapoor Amita, Oal Sujit; Deep Learning with Tensorflow 2.0 and Keras; 2019; Birmingham

Rudolf Böhm, Alsbach

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun in their paper – “Deep Residual Learning for Image Recognition”.

https://github.com/paddygoat/bioacoustics/blob/master/Train_and_deploy_bats/readme

https://en.wikipedia.org/wiki/Precision_and_recall

- i Wwww.tensorflow.org
- ii www.tierstimmenarchiv.de
- iii Rudolf Böhm, Alsbach-Hähnlein
- iv https://github.com/paddygoat/bioacoustics/blob/master/Train_and_deploy_bats/readme (house_keys, Junk_A ...
Junk_F)
- v <https://cran.r-project.org/web/packages/bioacoustics/index.html>
- vi https://en.wikipedia.org/wiki/Precision_and_recall